

Hw2

The task for Q1-Q3

- Class labels: three newsgroups
 - talk.politics.guns
 - talk.politics.mideast
 - talk.politics.misc
- Training data: 2700 instances (900 for each class)
- Test data: 300 instances (100 for each class)
- Features: words
- Task:
 - (Q1-Q2) Run Mallet DT learner
 - (Q3) Build your own DT learner

Use Mallet

- `mallet import-svmlight --input train.vectors.txt --output svmltrain.vectors`
 - Format of `train.vectors.txt`: `label f1:v1 f2:v2 ...`
 - `mallet train-classifier --input train.vectors --trainer MaxEnt --output-classifier m1`
 - Trains MaxEnt classifier and stores model
 - `mallet classify-svmlight --input test.vectors.txt --classifier m1 --output res`
 - Tests on the new data and the result is written to “res”
-
- `mallet import-svmlight --input test.vectors.txt --output test.vectors --use-pipe-from train.vectors`
 - `vectors2classify --training-file train.vectors --testing-file test.vectors --trainer DecisionTree --report test:raw test:accuracy test:confusion train:confusion train:accuracy > de1.stdout 2>de1.stderr`

Q3: build a DT learner

- Each node checks exactly one feature
- Features are all binary; that is, a feature is either present or non-present
 - ➔ The DT is a binary tree
- Quality measure: Information gain

Efficiency issue

- To select the best feature, you will need to calculate the info gain for each feature
- Therefore, you will need to calculate the counts of (c, f) and $(c, \text{not } f)$ for each class label c and each feature f .
- Try to do this efficiently.
- Report running time in Tables 2 and 3.

Patas usage

- When testing your code, use small data sets and small depth values first.
- Learn to use condor submit if your code runs for more than a couple of minutes.
- Always monitor your jobs.

Condor submit example

- For a command we can run as:
 `mycommand -a -n <mycommand.in >mycommand.out`
- The submit description file might look like this:

`Executable = mycommand`

`getenv = true`

`input = mycommand.in`

`output = mycommand.out`

`error = mycommand.error`

`Log = mycommand.log`

`arguments = "-a -n"`

`transfer_executable = false`

`request_memory = 2*1024`

`Queue`

Some useful Condor commands

- `condor_submit` — submit a job
- `condor_status` — list available nodes and their status
- `condor_q` — list the job queue
- `condor_rm` — remove a job from the queue

For more info about condor submit, please see

<https://wiki.ling.washington.edu/bin/view.cgi/Main/HowToUseCondor> and
<https://wiki.ling.washington.edu/bin/view.cgi/Main/CondorClusterHomepage>