

libSVM

LING572

Fei Xia

Documentation

- <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
- The libSVM directory on Patas:
 - /NLP_TOOLS/ml_tools/svm/libsvm/latest/
 - README
 - FAQ.html
 - svm-train, svm-predict, etc.
- More info:
 - [A practical guide to support vector classification](#)
 - [LIBSVM : a library for support vector machines](#)

Steps for using libSVM

- Define features in the input space (if use one of the pre-defined kernel functions)
- Scale the data before training/test
- Choose a kernel function
- Tune parameters using cross-validation

Main commands

- `svm-scale`: scaling the data (not needed for hw8)
- `svm-train`: training
- `svm-predict`: decoding

svm-train

- `svm-train [options] training_data model_file`
- Options:
 - t [0-3]: kernel type
 - g gamma: used in polynomial, RBF, sigmoid
 - d degree: used in polynomial
 - r coef0: used in polynomial, sigmoid
- Type “svm-train” to see options

Kernel functions

-t kernel_type : set type of kernel function (default 2)

0: linear: $u' * v$

1: polynomial: $(\text{gamma} * u' * v + \text{coef0})^{\text{degree}}$

2: RBF: $\exp(-\text{gamma} * |u - v|^2)$

3: sigmoid: $\tanh(\text{gamma} * u' * v + \text{coef0})$

svm-predict

- `svm-predict test_data model_file output_file`
- `svm-predict` produces only the system prediction in `output_file`.
- You will implement your own decoder in Hw8.

The format of training/test data

- Sparse format: no need to include features with value zero.

- Mallet format:

truelabel f1: v1 f2: v2

- libSVM format:

truelabel_idx feat_idx1:v1 feat_idx2:v2

(feat_idx, v) is sorted according to feat_idx in ascending order.

Ex: 1 20:1 23:0.5 34:-1 ...

When there are two classes

Classifying an instance x

$$\begin{aligned} f(x) &= \sum_i \alpha_i y_i K(x_i, x) - \rho \\ &= \sum_i \text{weight}_i K(x_i, x) - \rho \end{aligned}$$

where y_i (i.e., x_i 's label) is $+1$ ("c₀") or -1 ("c₁").

if $f(x) > 0$

then label it with c_0

else label it with c_1

Notation differences

	In SVM paper	In libSVM
Model	x_i, y_i, α_i b	$weight_i, x_i$ ρ
Prediction	$\sum_i \alpha_i y_i K(x_i, x) + b$	$\sum_i weight_i K(x_i, x) - \rho$
Representing y_i in training/test/output	+1 -1	0 1

System output of svm-predict

0 $##\ c_0$

0

1 $##\ c_1$

1

0

0

1

0

The format of the model file

svm_type c_svc

kernel_type rbf

gamma 0.5

nr_class 2

total_sv 535

rho 0.281122

label 0 1

nr_sv 272 263

SV

0.98836

...

This is weight for the support vector,
which is equal to $\alpha_i y_i$.

0:1 1:1 2:1 3:1 4:1 5:1 ...

This is a support vector with the format f1:v1 f2:v2 ...

Additional slides

Scaling the data

- To avoid features with larger variance to dominate those with smaller variance.
- Scale each feature to the range $[-1,+1]$ or $[0,1]$.
 - $[0,1]$ is faster than $[-1,1]$

svm-scale

- `svm-scale -l -1 -u 1 -s range_file training_data > training_data.scale`
- `svm-scale -r range_file test_data > test_data.scale`
- Scale feature values to $[-1, 1]$ or $[0,1]$
- No need to scale the data for Hw8.

When there are C classes

Handling a multi-class task

- All-pair
- Build a classifier for every (c_m, c_n) pairs
 - There are $C(C-1)/2$ classifiers
- The classifiers are stored in a compact format.

The format of the model file (when there are $C > 2$ classes)

```
svm_type c_svc
kernel_type rbf
gamma 0.5
nr_class 3
total_sv 2698
rho -0.0111642 -0.00216906 0.00951624
label 0 1 2
nr_sv 900 898 900
SV
0.98836 0.9975 0:1 1:1 2:1 3:1 4:1 5:1 ...
...
```

The rho array

It contains $C(C-1)/2$ elements, one per classifier

0 vs. 1, 0 vs. 2, ..., 0 vs. C-1,

1 vs. 2, 1 vs. 3, ..., 1 vs. C-1

2 vs. 3, ..., 2 vs. C-1

...

C-2 vs. C-1

The format of the SV line

Each line includes $C-1$ weights (i.e., $y_i \alpha_i$) followed by the vector.

$w_1 w_2 \dots w_{C-1} f1:v1 f2:v2 \dots$

Suppose the current vector belongs to the i -th class, the weights are ordered as follows:

0 vs. i 1 vs. i 2 vs. i $i-1$ vs. i
 i vs. $i+1$ i vs. $i+2$ i vs. $i+3$ i vs. $C-1$

Ex1: $i=0$

0 vs. 1, 0 vs. 2, 0 vs. 3, ..., 0 vs. $C-1$

Ex2: $i=4$

0 vs. 4, 1 vs. 4, 2 vs. 4, 3 vs. 4, 4 vs. 5, 4 vs. 6, ..., 4 vs. $C-1$

Classifying an instance x

$win[m]=0$ for every class m

For each classifier for (m,n)

$$\begin{aligned} f(x) &= \sum_i \alpha_i y_i K(x_i, x) - \rho \\ &= \sum_i \underbrace{\alpha_i y_i}_{weight_i} K(x_i, x) - \rho \end{aligned}$$

where x_i is a training instance with label c_m or c_n .

if $f(x) > 0$

then $win[m]++$

else $win[n]++$

$sysLabel = \arg \max_m win[m]$

To classify x with a m -vs- n classifier ($m < n$):

ρ is stored at what position?

For each x_i belonging to c_m

0 vs. m , 1 vs. m , ..., $m-1$ vs. m ,

m vs. $m+1$, m vs. $m+2$, ..., m vs. n , ...

the weight for m -vs- n is stored at position $n-1$

For each x_i belonging to c_n

0 vs. n , 1 vs. n , 2 vs. n , ..., m vs. n , ...

the weight for m -vs- n is stored at position m